

# Asynchronous Programming

Andrew Clymer  
andy@rocksolidknowledge.com

Richard Blewett  
richard@rocksolidknowledge.com



# Course prerequisites

- **Programming experience required**
  - 6 months of C#



# The plan

- **Tasks**
- **Async/Await**
- **Thread Safety**
- **Simplifying Thread Safety**



# Agenda


- **Why async programming**
- **The Task abstraction**
- **Creating Tasks**
- **Passing data into tasks and retrieving results**
- **Cancellation and Progress**

# The Benefits of Asynchronous Programming

- **Single threaded applications simplest to write**
  - Not practical for server side
  - Can cause responsiveness problems on client
- **Running code asynchronously has many potential benefits**
  - Increased throughput
  - Increased responsiveness
- **Async required to take advantage of modern CPU architectures**
  - Clock speeds not increasing significantly
  - Multi-core architecture now the norm



# The Free Lunch is Over



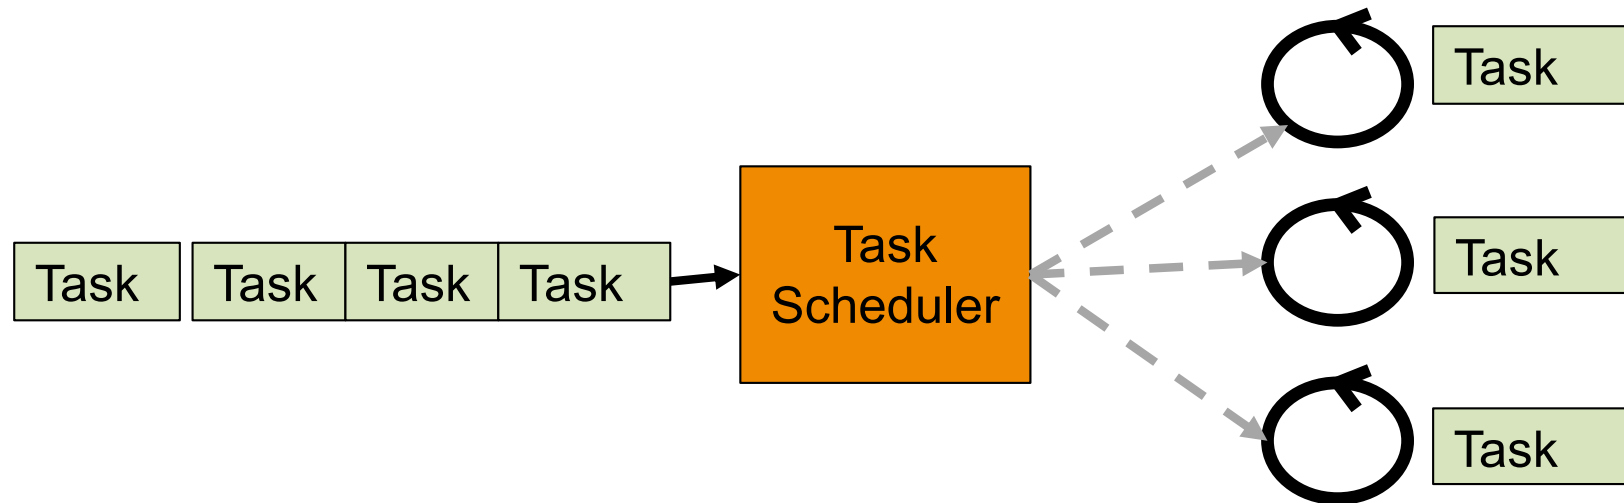
# Creating Asynchronous Units of Work

- **Developers indentify functionality that can run asynchronously**
  - Invoking “long running” blocking operations such as network access
  - Waiting for a condition to arise such as a file to be created in a particular folder
  - Processing a request from another machine
- **Work must be totally or mostly independent from main processing**
  - Otherwise coordinating work removes the benefit of async
- **Developer packages unit of work as a Task**



# What is a Task?

- A Task is a schedulable Unit of Work
  - Wraps a delegate that is the actual work
- Task is enqueued to a TaskScheduler





# Creating a Task

- Tasks are created by passing a **delegate to the constructor**
  - Call `Start` to queue the task to the scheduler
  - Can also use **Factory** property

```
Action a = () => Console.WriteLine("Hello from task");  
Task tsk = new Task(a);  
tsk.Start();
```

```
Action a = () => Console.WriteLine("Hello from task");  
Task tsk = Task.Factory.StartNew(a);
```




# Scheduling the Work

- **Scheduler maps work on to threads**
  - Scheduler is an abstraction that has multiple implementations
- **Two main models for scheduling**
  - Thread pooling
  - Dedicated threads
- **Thread pooling reuses existing threads**
  - Ideal for most situations
- **Dedicated threads give the task its own thread**
  - Is good practice for long running tasks

# The System Thread Pool

- .NET has a per-process thread pool
  - Assigned work queued up
  - Threads in pool pick up tasks when idle
- Thread pool manager controls number of threads in pool
  - Number of threads depends on workload
  - Number of threads in pool is capped



# A Single API

- Task provides a single API for async work
  - Long running task created by hinting to scheduler
  - Current implementation long running tasks spawn own thread otherwise Thread Pool used

```
Task t1 = new Task( DoWork,  
TaskCreationOptions.LongRunning );
```

# Passing Data to a Task

- Data passed explicitly using `Action<object>`
- Data can be passed **implicitly using anonymous delegate** rules

```
Guid jobId = Guid.NewGuid();

Action<object> a = state =>
{
    Console.WriteLine("{0}: Hello {1}", jobId, state);
};

Task tsk = new Task(a, "World");
tsk.Start();
```

# Returning Data from a Task

- **Generic version of Task available**
  - T is return type
  - Accessed from the task **Result**
- Takes a **Func** delegate as a constructor parameter
  - **Func<T>**
  - **Func<object, T>**

```
Func<object, int> f = state =>
{
    Console.WriteLine("Hello {0}", state);
    return 42;
};
Task<int> tsk = new Task<int>(f, "World");
tsk.Start();
//...
Console.WriteLine("Task return is: {0}", tsk.Result);
```

# Waiting for a Task to End

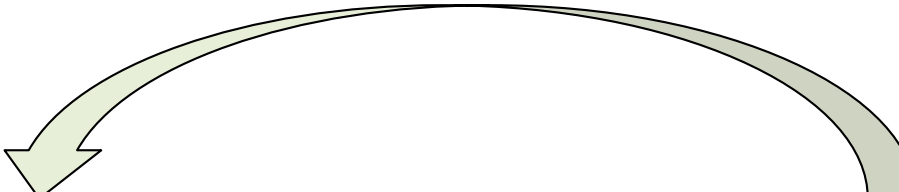
- Can wait for one or more tasks to end using `Wait`, `WaitAll` or `WaitAny`
- Can pass **timeout** for wait

```
Task t = new Task( DoWork );  
t.Start();  
t.Wait();
```

```
Task t1 = new Task( DoWork );  
t1.Start();  
Task t2 = new Task( DoOtherWork );  
t2.Start();  
  
if ( !Task.WaitAll( new Task[] { t1, t2 }, 2000 ) )  
{  
    Console.WriteLine( "wait timed out" );  
}
```

# Task Completion States

- Tasks can end in one of three states
  - RanToCompletion: everything completed normally
  - Canceled: task was cancelled
  - Faulted: an unhandled exception occurred on the task
- Unhandled Exceptions get thrown when **waiting on a task**
  - thrown in the Task's finalizer if task not waited upon



```
Task t = new Task(DoWork);  
t.Start();  
  
if (!t.Wait(1000))  
{  
}
```

```
private static void DoWork()  
{  
    throw new Exception();  
}
```



# Cancellation

- **Tasks support cancellation**
  - Modelled by `CancellationToken`
- **Token can be passed into many APIs**
  - Task creation
  - **Waiting**

```
var source = new CancellationTokenSource();  
  
Task t1 = new Task( DoWork, source.Token );  
t1.Start();  
  
t1.Wait(source.Token);
```

# Triggering Cancellation

- **CancellationTokenSource** has **Cancel** method to trigger the cancellation of tasks and blocking APIs

```
var source = new CancellationTokenSource();  
  
Task t1 = new Task( DoWork, source.Token );  
t1.Start();  
  
source.Cancel();
```

# The Effects of Cancellation

- Cancellation has different effects depending on state of task
  - Unscheduled tasks are never run
  - Scheduled tasks must cooperate to end. **Requires access to CancellationToken**

```
private static void DoWork(object o)
{
    var tok = (CancellationToken)o;

    while (true)
    {
        Console.WriteLine("Working ...");
        Thread.Sleep(1000);
        tok.ThrowIfCancellationRequested();
    }
}
```

# Cancellation of Blocking Operations

- **Blocking operations throw exceptions when cancelled**
  - **OperationCancelledException**
  - **AggregateException** (when more than one exception could have occurred)

```
try
{
    t1.Wait(1000, source.Token);
}
catch (AggregateException x)
{
    foreach (var item in x.Flatten().InnerExceptions)
    {
        Console.WriteLine(item.Message);
    }
}
```

# Timeout via cancellation

- Can signal **cancellation** after a given period of time

```
var source = new CancellationTokenSource();  
  
source.CancelAfter(TimeSpan.FromSeconds(2));  
  
Task t1 = new Task( DoWork, source.Token );  
t1.Start();
```

# Linking Cancellation

- **Task** cancelled when either
  - **ctsToken** is signaled
  - **timeoutCancel** source signals

```
static Task DoItAsync( in CancellationToken ctsToken) {  
  
    var timeoutCancel = CancellationTokenSource  
        .CreateLinkedTokenSource(ctsToken);  
  
    timeoutCancel.CancelAfter(TimeSpan.FromSeconds(3));  
  
    var result = ExecuteCallAsync(timeoutCancel.Token);  
  
    return result;  
}
```

# Task Completion Source

- Tasks can be used to represent any asynchronous activity
- `TaskCompletionSource<T>` used to manage own Task lifetime
- Used to adapt non Task based asynchronous operations to Tasks

# Wrapping an event

- Single raised event wrapped as a task


```
public static Task<FileSystemEventArgs> CompleteOnChange(  
    this FileSystemWatcher watcher )  
{  
    var tcs = new TaskCompletionSource<FileSystemEventArgs>();  
  
    void OnWatcherOnChanged(object sender, FileSystemEventArgs args)  
    {  
        watcher.Changed -= OnWatcherOnChanged;  
        tcs.SetResult(args);  
    }  
    watcher.Changed += OnWatcherOnChanged;  
    return tcs.Task;  
}
```



# Async IO

- **Some .NET APIs model async IO**
  - No thread is consumed while IO bound operation takes place
  - Uses IO Completion ports
- **Async IO modelled with Async method**
  - Long running operation has additional method that returns a `Task<T>`

```
WebResponse GetResponse ()
```



```
Task<WebResponse> GetResponseAsync ()
```

# Scheduling

- TaskScheduler is an extensible abstract class
- Out of the box implementations
  - ThreadPoolTaskScheduler (Default)
  - SynchronizationContextTaskScheduler
- Can pass scheduler when starting a task

```
TaskScheduler scheduler = GetScheduler();  
  
Task t = new Task(DoWork);  
  
t.Start(scheduler);
```

# Integrating with SynchronizationContext

- GUI applications need UI updates marshalled to the UI thread
  - SynchronizationContext is abstraction that wraps specific technology solution to thread marshalling
- Can get scheduler associated with SynchronizationContext

On UI Thread

```
TaskScheduler scheduler =  
    TaskScheduler.FromCurrentSynchronizationContext();
```

On Background Thread

```
private void DoAsyncWork()  
{  
    Task t = new Task(() => label1.Text = "TADA!!");  
    t.Start(scheduler);  
}
```

# Unnecessary task allocation

- **Synchronous result** still result in a GC allocation

```
public class StockDataSource {  
    public Task<decimal> GetPrice(string symbol) {  
        if (localPrices.TryGetValue(symbol, out decimal price)){  
            return Task.FromResult(price);  
        } else {  
            return GetRemotePrice(symbol);  
        }  
    }  
}
```

# ValueTask<T>

- Value Type
- No allocation for completed tasks
  - Reduce load on GC
- Designed for high throughput scenarios where operations often complete synchronously
- Useful for abstractions that may have synchronous or asynchronous implementations

# Removing unnecessary task allocation

- **Synchronous result** will not result in a GC allocation
- **Asynchronous result** wrapped by a ValueTask

```
public class StockDataSource {  
    public ValueTask<decimal> GetPrice(string symbol) {  
        if (localPrices.TryGetValue(symbol, out decimal price)){  
            return new ValueTask<decimal>(price);  
        } else {  
            return new ValueTask<decimal>(GetRemotePrice(symbol));  
        }  
  
        private Task<decimal> GetRemotePrice(string symbol) { ... }  
    }  
}
```

# Should all asynchronous APIs be `ValueTask<T>`?


- `Task<T>` results can be safely used multiple times and concurrently
- `ValueTask<T>`, can not be safely used multiple times and concurrently
- `ValueTask<T>` not as feature rich as `Task<T>`
  - No `Task.WhenAny`, `Task.WhenAll`, `Task.WaitAll` ...
  - Can convert to task using `.AsTask()`
- Simple answer is no

# Summary

- Async programming is an important skill
- Tasks model unit of async work
  - Compute
  - IO
- Tasks support cancellation
- TaskCompletionSource useful tool for making anything a task
- Consider using ValueTask for abstractions where the implementation may often be synchronous



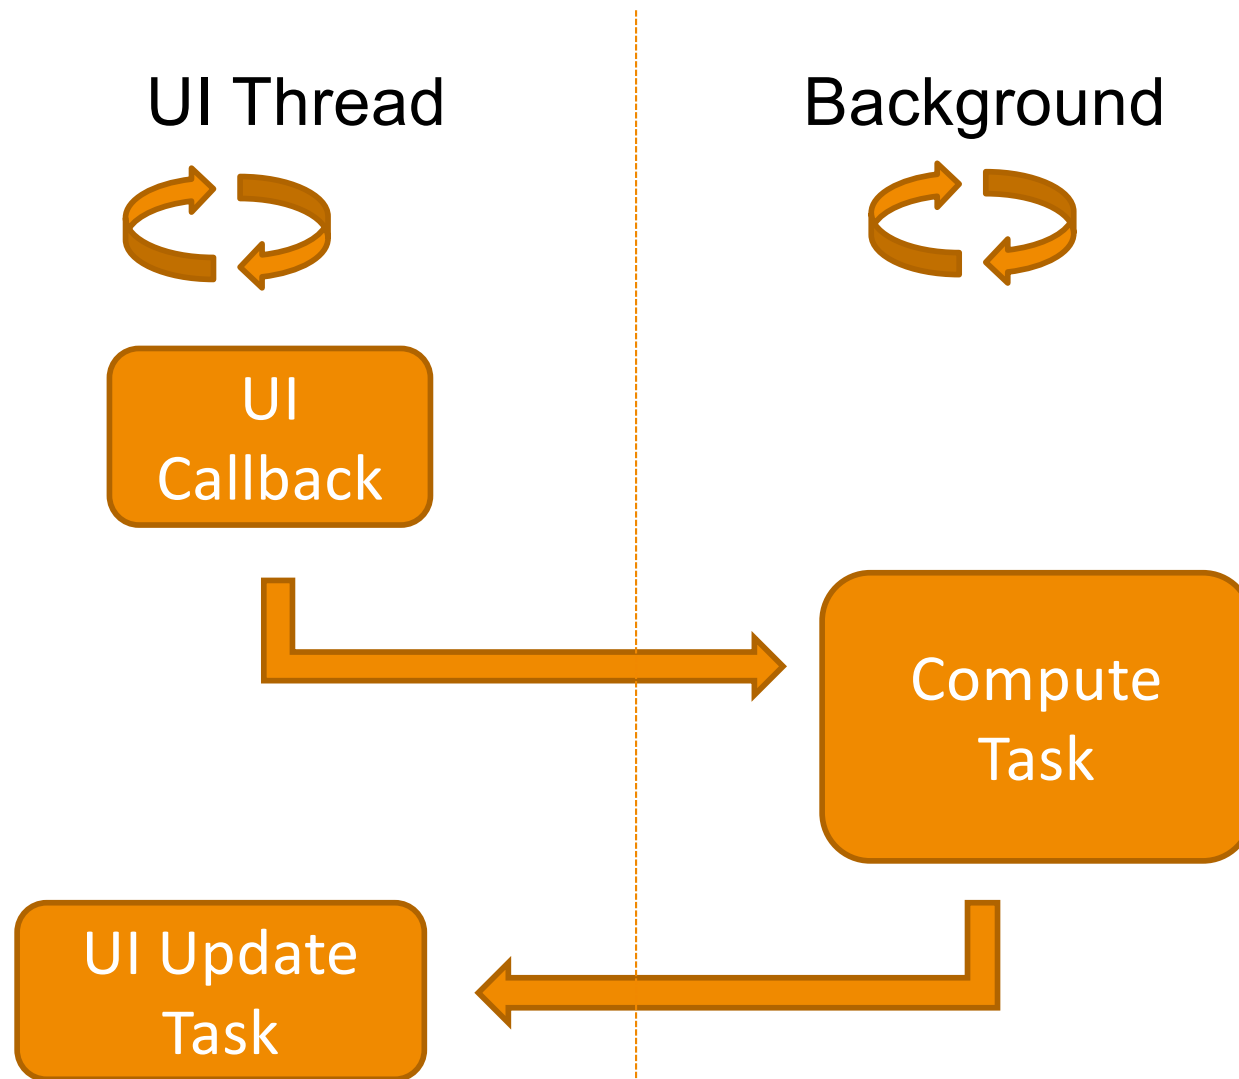
# async and await



# Objectives

- **Why use Continuations**
- **Simple examples of async/await**
- **Under the hood**
- **Gotcha's**
- **Composition**
- **Server side async/await**
- **Asynchronous streams**

# Continuations



# Control flow

- Sequential programming intent is pretty clear.
- Asynchronous programming screws with intent
- Do X Async, Then Do Y , Then Do Z Async
- How to handle errors
  - Where to place the try/catch

# async/await keywords

Making async intent as clear as synchronous intent

- Two new keywords introduced in C# 5
- Enables continuations whilst maintaining the readability of sequential code
  - Automatic marshalling back on to the UI thread
- Built around Task, and Task<T>, ValueTask<T>

# async and await

## Example

- **async** method must return void or a Task/ValueTask
- **async** method should include an **await**
- **await** <TASK>

```
async void Button_Click(object sender, RoutedEventArgs e) {  
    calcButton.IsEnabled = false;  
    Task<double> piResult = CalcPiAsync(1000000000);  
  
    // If piResult not ready returns, allowing UI to continue  
    // When has completed, returns back to this thread  
    // and coerces piResult.Result out  
        await piResult;  
  
    calcButton.IsEnabled = true;  
    this.pi.Text = pi.ToString();  
}
```

# async await

Can return Task<T>

- Code returns **T**, compiler returns **Task<T>**

```
async Task<byte[]> DownloadDataAsync(Uri source)
{
    WebClient client = new WebClient();
    byte[] data = await client.DownloadDataTaskAsync(source);

    ProcessData(data);

    return data
}
```

# Favour continuations over waiting

- Threads aren't free
- A thread waiting can't be used for anything else.
- Using continuations can reduce the total number of required threads



# async/await under the hood

## Compiler builds state machine

```
private static async void TickTockAsync()  
{  
    Console.WriteLine("Starting Clock");  
  
    while (true)  
    {  
        Console.WriteLine("Tick");  
        await Task.Delay(500);  
  
        Console.WriteLine("Tock");  
        await Task.Delay(500);  
    }  
}
```

Console.WriteLine("Starting  
Clock");

Console.WriteLine("Tick");  
await Task.Delay(500);

Console.WriteLine("Tock");  
await Task.Delay(500);

**async** keyword does not make code run asynchronously

```
async Task DoItAsync()
{
    // Still on calling thread
    Thread.Sleep(5000);
    Console.WriteLine("done it..");
}
```

### Avoid **async** methods returning **void**

```
async void DownloadDataAsync(string uri){  
    . . .  
}
```

- Better to return **Task** than void
- Allows caller to handle error
- void is there for asynchronous event handlers

```
async Task DownloadData(Uri source)  
{  
    WebClient client =new WebClient();  
    byte[] data = await client.DownloadDataTaskAsync(source);  
  
    ProcessData(data);  
}
```

### THINK before using **async** lambda for Action delegate

```
requests.ForEach(async request =>
{
    var client = new WebClient();
    Console.WriteLine("Downloading {0}", request.Uri);
    request.Content = await
        client.DownloadDataTaskAsync(request.Uri);
});
Console.WriteLine("All done...??");

requests.ForEach(r => Console.WriteLine(r.Content.Length);
```

# async/await

## Gotcha #4

- **await exception handling only delivers first exception**
- **Tasks can throw many exceptions via an `AggregateException`**
  - Await re-throws only **first exception** from Aggregate
- **Examine `Task.Exception` property for all errors**

```
Task<byte[]> loadDataTask = null;
try {
    loadDataTask = LoadAsync();
    byte[] data = await loadDataTask;
    ProcessData(data)
} catch (Exception firstError) {
    loadDataTask.Exception.Flatten().Handle( MyErrorHandler );
}
```

# ConfigureAwait

Possibly the worst API ever conceived

- Not all **await's** need to make use of **SynchronizationContext**

```
public static async Task DownloadData(Uri source, string destination)
{
    WebClient client = new WebClient();
    byte[] data = await client.DownloadDataTaskAsync(source);

    // DON'T NEED TO BE ON UI THREAD HERE...
    ProcessData(data);

    using (Stream downloadStream = File.OpenWrite(destination))
    {
        await downloadStream.WriteAsync(data, 0, data.Length);
    }
    // Must be back on UI thread
    UpdateUI("Download
```

# ConfigureAwait

Possibly the worst API ever conceived

- **First attempt, but wrong**

```
static async Task DownloadData(Uri source, string destination)
{
    WebClient client = new WebClient();
        await client
            .DownloadDataTaskAsync(source)
            .ConfigureAwait(false);

    // Will continue not on UI thread
    Stream downloadStream = File.OpenWrite(destination) {
        await downloadStream.WriteAsync(data, 0, data.Length);
    }

    // Hmmm...Need to be back on UI thread here

    UpdateUI("All downloaded");
}
```

# ConfigureAwait

## Effective use of ConfigureAwait with composition

- Get compiler to create **Task** per context

```
static async Task DownloadData(Uri source, string destination){
    await DownloadAsync(source, destination);
    // on UI thread
    UpdateUI("All downloaded");
}

private static async Task DownloadAsync(Uri source, string destination) {
    WebClient client = new WebClient();
    byte[] data = await client
        .DownloadDataTaskAsync(source)
        .ConfigureAwait(continueOnCapturedContext:false);

    using (Stream downloadStream = File.OpenWrite(destination)) {
        await downloadStream.WriteAsync(data, 0, data.Length);
    }
}
```



# Await and Monitors

- Compile will not allow an **await** inside a **lock** block
  - Lock only works if the entire block is executed on the same thread

```
lock (account)
{
    await UpdateAccount().ConfigureAwait(false);
}
```

Won't compile

```
Monitor.Enter(account)
try
{
    await UpdateAccount().ConfigureAwait(false);
}
finally{ Monitor.Exit(account); }
```

Will compile

# Utilise Semaphore slim

## Await compatible locking

- Semaphore with count of 1 has similar behavior to that of Mutex
- Semaphore can be acquired and released around a await
- Can await on a semaphore, for non blocking synchronization
- Wrap in using pattern to maintain programming model

# Awaiting with timeout

- **Waiting for ever, is bad**
- **Awaiting for ever possibly less bad**
- **What if awaiting task has no cancellation?**
  - **await keyword has no time out**
- **Consider using Task.WhenAny**
- **.NET 6 introduces Task.WaitAsync()**

# Async on the server

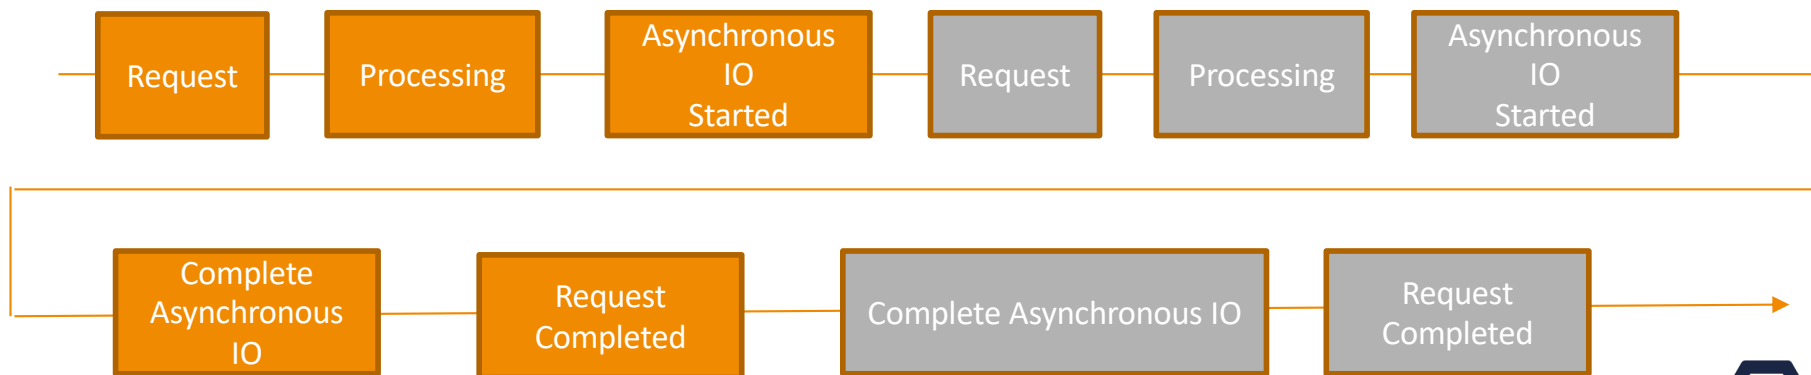
Not just client side technology

- **MVC and WebAPI both understand**
  - `Task<T>`
  - `ValueTask<T>`
- **Blocking on a thread is harmful to your application**
  - New requests may force a new thread to be created
  - New threads take time to start and consume resource
- **Server threads shouldn't block**
  - Release thread until they need it again
  - Allow N concurrent requests to share a single thread of execution

# Async on the server

## Thread re-use

- Thread starts processing a given **request**
- Initiates some asynchronous IO, and frees itself to perform another **request**
- Second request initiates asynchronous IO, and frees itself to complete the previous **request**



# AsyncLocal<T>

- Applications sometimes need to flow ambient state
  - HttpContext
  - LoggingContext
- When using async/await can't use thread local storage
- AsyncLocal<T> used to flow ambient state across await boundaries

# Async Enumerable aka Async Streams

- Why the need for asynchronous iteration
- Foreach async
- Disposable Async

# DIY async iteration

- Asynchronous iteration awkward for creation and consumption

```
public static IEnumerable<Task<string[]>> LoadCsv(string filename){  
    using (var reader = new StreamReader(filename)) {  
        while (!reader.EndOfStream){  
            yield return LoadAndSplit(reader);  
        }  
    }  
}  
  
private static async Task<string[]> LoadAndSplit(StreamReader reader) {  
    return (await reader.ReadLineAsync()).Split(',');  
}
```



# IAsyncEnumerable<T>

- Async version of IEnumerable, IEnumerator

```
public interface IAsyncEnumerable<out T> {  
    IAsyncEnumerator<T> GetAsyncEnumerator(Cancellation token ct);  
}  
  
public interface IAsyncEnumerator<out T> : IAsyncDisposable {  
  
    T Current {get; }  
    ValueTask<bool> MoveNextAsync();  
}
```

# Yield return async enumerable

- Compiler creates an async enumerable
- Methods must be marked **async**
- Methods must return **IAsyncEnumerable<T>** or **IAsyncEnumerator<T>**
- **Yield return** used as per iterator methods

```
static async IAsyncEnumerable<string[]> LoadCsv(string filename) {  
    using (var reader = new StreamReader(filename))  
    {  
        while (!reader.EndOfStream) {  
            string row = await reader.ReadLineAsync();  
            yield return row.Split(',');  
        }  
    }  
}
```

# Foreach async

- Iterates through **async enumerable**
- Delivers each item, not Task<T>

```
IAsyncEnumerable<string[]> rows = LoadCsv(@"stockData.csv");  
  
await foreach (string[] row in rows)  
{  
    Console.WriteLine(row[1]);  
}
```

# Async LINQ

- Nuget package System.Linq.Async
- `IAsyncEnumerable<T>` extension methods defined in **AsyncEnumerable**

```
var tradingDays = LoadCsv(@"stockData.csv")
    .Skip(1)
    .Select(row => new
    {
        When = DateTime.Parse(row[0]),
        Open = decimal.Parse(row[1]),
        Close = decimal.Parse(row[4])
    });

await foreach (var tradingDay in tradingDays) {
    Console.WriteLine(tradingDay);
}
```

# Foreach async cancellation

- Iterator method can take cancellation token
- Problem: this is scoped for the `IAsyncEnumerable` not the `IAsyncEnumerator`

```
var cts = new CancellationTokenSource();
var rows = LoadCsv("stockData.csv" , cts.Token);

await foreach (string[] row in rows) {
    Console.WriteLine(row[1]);
    cts.Cancel();
}

await foreach (string[] row in rows) {
    Console.WriteLine(row[1]); // NEVER EXECUTES
}

public static async IAsyncEnumerable<string[]> LoadCsv(
    string filename , CancellationToken ct)
```

# Enumeration cancellation

- Cancellation can be **scoped to the enumerator**
- Iterator method **parameter** marked to accept cancellation token

```
var cts = new CancellationTokenSource();
var rows = LoadCsv("stockData.csv" , CancellationToken.None);

await foreach (string[] row in rows.WithCancellation(ct.Token) {
    Console.WriteLine(row[1]);
    cts.Cancel();
}

public static async IAsyncEnumerable<string[]> LoadCsv(
    string filename ,
    [EnumeratorCancellation]CancellationToken ct)
```

# Async Disposable

- New interface **IAsyncDisposable**
- Using statement prefix with await uses **DisposeAsync**

```
public interface IAsyncDisposable {  
    ValueTask DisposeAsync();  
}
```

```
await using (FileStream s = File.OpenRead("stockData.csv")){  
  
}
```

# Async iteration on the server

- MVC and WebAPI both understand **`IAsyncEnumerable<T>`**

```
[Route("TradingDays")]
[HttpGet]
public IAsyncEnumerable<TradingDayDTO> GetStocks()
{
    var rows = Context.TradingDays;
    return rows.Select(r => new TradingDayDTO()
    {
        When = r.When.ToShortDateString(),
        Close = r.Close,
        Open = r.Open,
        Volume = r.Volume
    }).AsAsyncEnumerable();
}
```




# Summary

- **Utilise async/await to**
  - **Simplify continuations**
  - **Reduce number of threads**
- **Use Semaphore as an await safe locking primitive**
- **Use ConfigureAwait to reduce work on UI thread**
- **Use asynchronous apis for greater scalability and performance**
- **Use IAsyncEnumerable for streaming asynchronous results**



# Thread Safety



# Agenda

- Highlight issues with multi threaded programming
- Introduce thread synchronization primitives
- Introduce thread safe collections

# Need for Synchronization

- **Creating threads is easy**
- **When threads share data problems can occur**
  - Inconsistent reads
  - State corruption
- **Synchronization fixes these problems, but potentially creates a new problem**
  - Over synchronization reduces scalability
- **Lots of techniques to implement synchronization**
  - Each have cost and benefit
- **Developers role is to write an application that scales and is thread safe, by selecting the best synchronization technique**

# Simple Increment

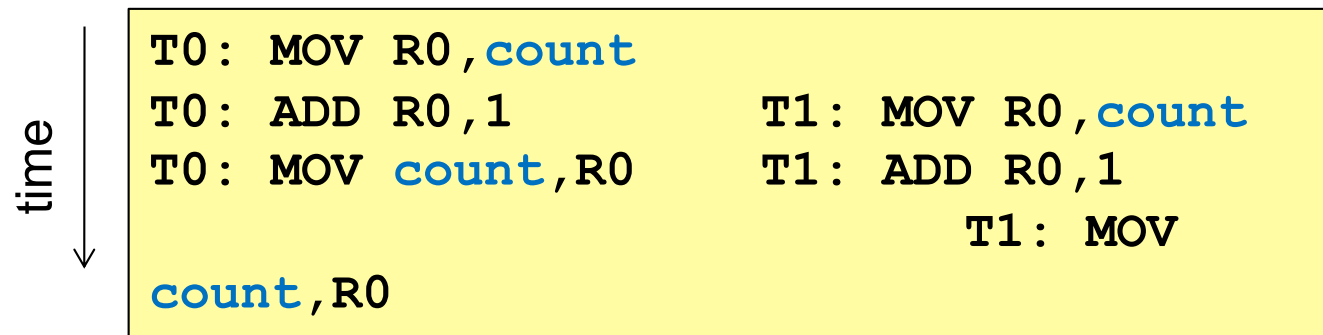
- Two threads
  - Sharing an instance of **Counter**.
  - Both are calling **Increment** 1000 times
- Question
  - What is the value of count after both threads have completed?

```
public class Counter
{
    protected int count;

    public virtual void Increment()
    {
        count++;
    }
    public int Value { get { return count; } }
}
```

# Simple Increment, NOT Atomic

- Even a simple `count++` is not an atomic operation.
  - Multiple CPU instructions that could be interweaved.



- Consider the possible execution below of two threads (T0, T1)
  - Assuming `count=0` at the start
  - At the end of execution i would be 1 and not the desired 2.
- If two threads don't attempt to increment count at the same time not a problem. Spotting these kind of errors is hard

# Interlocked

- Modern CPU's expose special instruction set to perform various operations atomically
  - Cost more than non atomic variants.

```
public class InterlockedCounter : Counter
{
    public override void Increment()
    {
        // Atomic count++
        Interlocked.Increment(ref count);
    }
}
```

- Access to these instructions via **Interlocked** class
  - Interlocked.Increment
  - Interlocked.Decrement
  - Interlocked.Add
  - Interlocked.Exchange, Interlocked.CompareAndExchange
    - Useful for building Spin locks

# Multi step state transition

- What happens if
  - Thread A is inside ReceivePayment
  - Thread B is **inside NetWorth**
- Can Interlocked help ?

```
class SmallBusiness {  
    decimal Cash = 0;  
    decimal Receivables = 1000;  
  
    public void ReceivePayment(decimal amount){  
        Cash += amount;  
        Receivables -= amount;  
    }  
  
    public decimal NetWorth {  
        get { return Cash + Receivables; }  
    }  
}
```



# Sequential access

- To fix the problem
  - Sequentialise access to the object state
- How
  - Each instance of a reference type has a Monitor
  - CLR guarantees that only one thread can own the monitor
  - If a thread can't acquire the monitor it enters a wait state
  - When the monitor is available it is woken up and proceeds
- Critical areas of code can therefore be protected by using a monitor.

# Monitor based solution

- Only one thread in any **critical region** at any point in time

```
private object _lock = new object();

public void ReceivePayment(decimal amount)
{
    Monitor.Enter(_lock);
    Cash += amount;
    Receivables -= amount;
    Monitor.Exit(_lock);
}
```

Could be an issue with exceptions

```
public decimal NetWorth
{
    get {
        Monitor.Enter(_lock);
        try { return Cash + Receivables; }
        finally { Monitor.Exit(_lock); }
    }
}
```

Deals better with exceptions

# Lock keyword

- Enter, try, finally , Exit common pattern
  - C# language offers **lock** keyword to assist
  - Compiler emits try, finally logic
- Use of Monitor.Enter and **lock** can lead to deadlocks
  - Prefer Montior.TryEnter which takes a timeout
- Avoid using lock[this] and lock[typeof(X)]
  - Less control over objects use for synchronization.
  - Prefer creation of object for sole purpose of

```
public void ReceivePayment(decimal amount){  
    lock (_lock)  
    {  
        Cash += amount;  
        Receivables -= amount;  
    }  
}
```

# High Read to Write Ratio

- Monitor provides mutual exclusion behaviour
  - Excluding readers and writers
- Thread safety not an issue if all threads read.
- Better throughput may be achieved with a Synchronization primitive that ensures
  - There can be Many Readers, Zero Writer
  - Or One Writer, Zero Readers
- This is known as a ReaderWriterLock
  - .NET 3.5 and above prefer ReaderWriterLockSlim
  - Pre 3.5, ReaderWriterLock
    - Not well implemented, can result in writer being denied access for long periods of time.
- Often used for caching



# Reader Writer Lock

```
public class SimpleCache
{
    private Dictionary<int, string> cache=
        new Dictionary<int, string>();
    private ReaderWriterLockSlim _lock = new ReaderWriterLockSlim();
```

```
    public string Get(int key) {
        _lock.EnterReadLock();
        try { return cache[key]; }
        finally { _lock.ExitReadLock(); }
    }
```

Many threads can  
can read from the  
cache

```
    public void Set(int key, string val)
    {
        _lock.EnterWriteLock();
        try { cache.Add(key, val) }
        finally { _lock.ExitWriteLock(); }
    }
}
```

When one thread  
has the write lock  
no other thread  
can obtain read or  
write lock.

# Semaphores

- Primitive that can be acquired for a set number of times
  - Often used to implement throttling
- Semaphore can be released from any thread, unlike Monitor
  - Semaphore count of one similar to semantics to Monitor
- Two implementations in the framework
  - SemaphoreSlim
  - Semaphore, wraps Kernel level primitive

# Synchronization across process

- Managed synchronization primitives only allow synchronization inside a single process
- How to control access to a shared file ?
  - Requires Kernel based synchronization
- Kernel synchronization can be achieved via managed wrappers
  - Mutex
  - Semaphore
  - AutoResetEvent
  - ManualResetEvent
- These synchronization primitives are orders of magnitude more expensive than managed ones




# Summary

- A variety of ways to perform synchronization, the skill is picking the correct one
- Concurrent collections make it simpler to write efficient thread safe code
- Only use kernel synchronization primitives when absolutely necessary
- Analyse code and imagine worse possible race conditions



# Simplifying thread safety



# Agenda

- **Greater concurrency without complexity**
- **Lazy<T>**
- **Concurrent Collections**
  - **Blocking Collections**
- **Channels .NET Core**
- **Immutable Collections**

# Thread Safe code can be complex

- Single threaded algorithms if written well simple to understand
- Add threads possibly require locking
  - Simple lock can be inefficient
  - Perhaps ReaderWriter Locks
  - Perhaps double check locking
- **RESULT: Complex code hard to maintain original intent often lost**

# Need for locks

- **Mutable shared state**
- **Shared state often takes the form of**
  - Collections [ List, Dictionary , Queue , Stack ]
- **Solutions**
  - Thread safe collections, hide synchronization
    - Concurrent collections
  - All shared state is immutable
    - Immutable collections

# Lazy<T>

- Provides thread safe on first read initialisation
  - Cheap stand in
- Useful
  - For delay loading the contents of a collection
  - Thread safe Virtual Proxy


# Concurrent Collections

- **Collections are the bedrock of most apps**
  - List, Dictionary, Queue, Stack
- **Problem, not thread safe**
- **Synchronized proxies/wrappers don't cut it**


# Consider this

```
Queue<int> queue = new Queue<int>();  
queue.Enqueue(1);  
. . .
```

```
if (queue.Count > 0)  
{  
    int val = queue.Dequeue();  
}
```



```
if (queue.Count > 0)  
{  
    int val = queue.Dequeue();  
}
```



# Concurrent Collections

## Concurrent API

- If/do, introduces race conditions
- Concurrent collection API remove if/do
  - TryXXXX
  - More complex atomic operations
    - AddOrUpdate
    - GetOrAdd
- **WARNING**...Be careful when using extension methods based on non-concurrent interfaces.
  - ToList()



# ConcurrentDictionary<K,V>

- **30-40% insert speed improvement in 4.5**
  - Re-use Nodes for reference and small value types
  - Number of locks change as structure grows
- **Initialise with potential size and level of concurrency for best performance**



# ConcurrentBag<T>

- List keeps items in order
- Bag keeps items
- What is it NOT
  - IT IS NOT A THREAD SAFE UNORDERED LIST
- It is ideally for load balancing divide/conquer

# What if I need to block

- Concurrent data structures don't block
  - Highly concurrent
- If require value before proceeding consider blocking
- **BlockingCollection<T>**
  - Adds block semantics to implementors of
    - **IProducerConsumerCollection <T>**

# Issue with blocking collections

- **Blocking a thread pool thread is RUDE**
- **To scale well**
  - Minimum number of threads maximum concurrency
- **async/await provides convent programming model to release and resume thread usage**

# Asynchronous queue

## .NET Core

- **Channel<T>** for producer consumer pattern
  - Supports asynchronous reads and writes
  - Support `IAsyncEnumerable`
- **Supports bounded and unbounded queues**
  - `Channel.CreateUnbounded<T>();`
  - `Channel.CreateBounded<T>(size);`
- **Can be optimized for**
  - Single Reader
  - Single Writer
  - Synchronous writes and reads

# Channel Types


- **Unbounded channel**
  - Assumes memory never runs out
  - Best for performance
- **Bounded channel**
  - Constrain number of items in the channel
  - Configurable when full behavior
    - Wait
    - Drop Newest
    - Drop Oldest
    - Drop Write

# Immutable Collections

- Thread Safety can be hard with mutable data It's a breeze with immutable data
- Not easy to achieve
- NuGet Microsoft Immutable collections
- Mutable operations results in efficient creation of new collection




# Add item D to Immutable list





# Remove item B to Immutable list



# Summary

- **Thread safety now achieved with high level abstraction**
  - Maintains readability
  - Greater confidence it works
  - Leverage on going development

