


Parallel Patterns




Objectives

- Why should I care about parallelism?
- Types of parallelism
- Introduction to Parallel Patterns



The free lunch is over



- Herb Sutter, Dr Dobb's Journal 30th March 2005

The continual need for speed

- For a given time frame X
 - Data volume increases
 - More detail in games
 - Intellisense
 - Better answers
 - Speech recognition
 - Route planning
 - Word gets better at grammar checking



Expectations

- Amdahl's Law
 - If only $n\%$ of the compute can be parallelised at best you can get an $n\%$ reduction in time taken.
 - Example
 - 10 hours of compute of which 1 hour cannot be parallelised.
 - So minimum time to run is 1 hour + Time taken for parallel part, thus you can only ever achieve a maximum of 10x speed up.
- Not even the most perfect parallel algorithms scale linearly
 - Scheduling overhead
 - Combining results overhead



Coarse grain parallelism

- Large tasks, little scheduling or communication overhead
 - Each service request executed in parallel, parallelism realised by many independent requests.
 - DoX,DoY,DoZ and combine results. X,Y,Z large independent operations
- Applications
 - Web Page rendering
 - Servicing multiple web requests
 - Calculating the payroll



Fine grain parallelism

- Fine grain parallelism
 - A single activity broken down into a series of small co-operating parallel tasks.
 - $f(n=0...X)$ can possibly be performed by many tasks in parallel for given values of n
 - Input \Rightarrow Step 1 \Rightarrow Step 2 \Rightarrow Step 3 \Rightarrow Output
 - Applications
 - Sensor processing
 - Route planning
 - Graphical layout
 - Fractal generation [Trade Show Demo]



Parallel Patterns

- Fine grain parallelism
 - A single activity broken down into a series of small co-operating parallel tasks.
 - $f(n=0...X)$ can possibly be performed by many tasks in parallel for given values of n
 - Input \Rightarrow Step 1 \Rightarrow Step 2 \Rightarrow Step 3 \Rightarrow Output
 - Applications
 - Sensor processing
 - Route planning
 - Graphical layout
 - Fractal generation [Trade Show Demo]



Fork and join

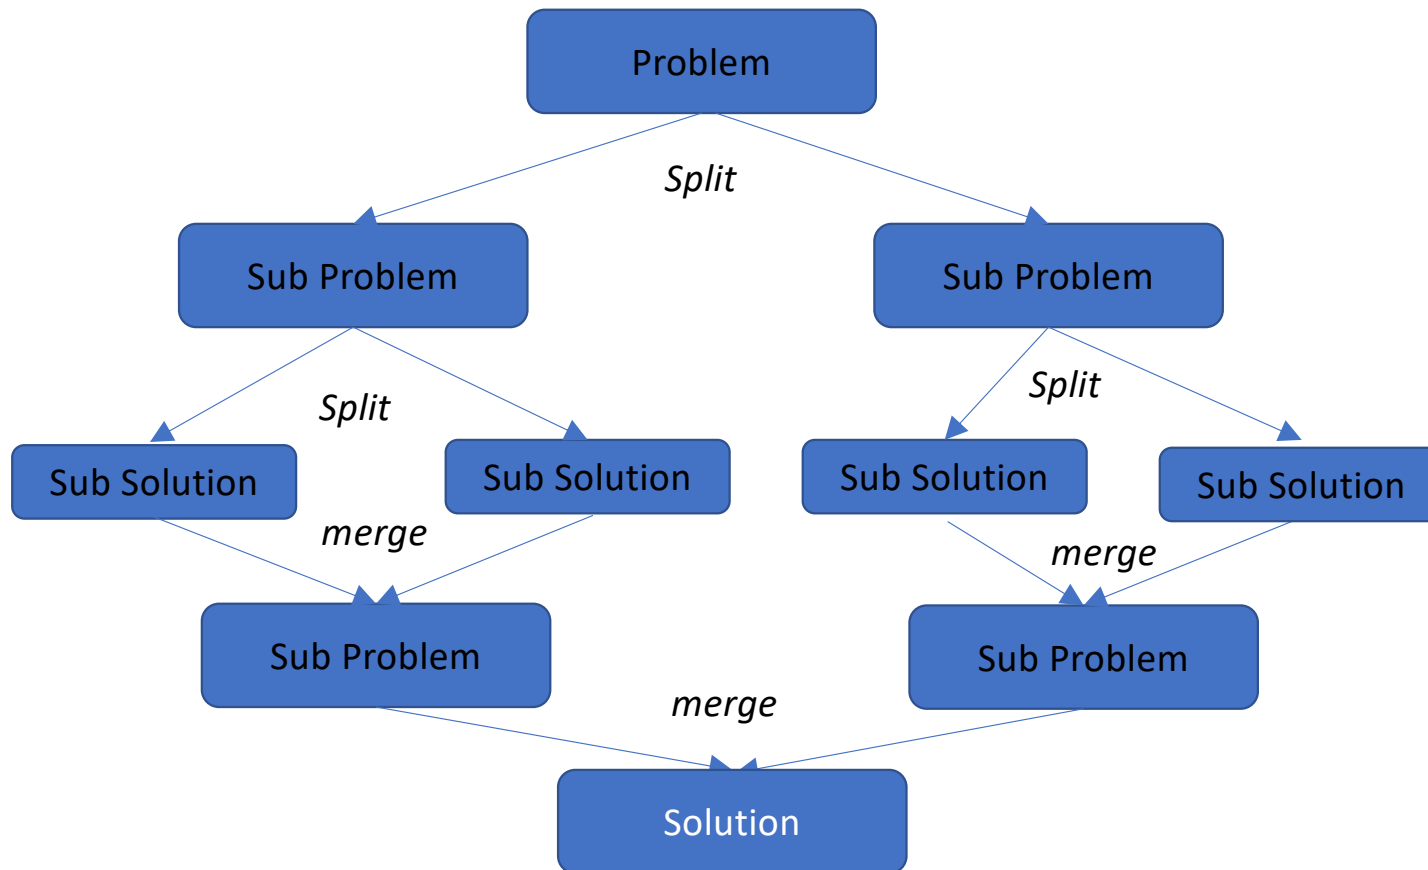
- Book Hotel, Book Car, Book Flight
- Sum X + Sum Y + Sum Z

```
Parallel.Invoke( BookHotel , BookCar , BookFlight );  
  
// Won't get here until all three tasks are complete
```

```
Task<int> sumX = Task.Factory.StartNew(SumX);  
Task<int> sumY = Task.Factory.StartNew(SumY);  
Task<int> sumZ = Task.Factory.StartNew(SumZ);  
  
int total = sumX.Result+ sumY.Result+ sumZ.Result;
```




Divide and Conquer



Geometric decomposition


- Task parallelism is one approach, data parallelism is another
- Decompose the data into a smaller sub set
- Have identical tasks work on smaller data sets.
- Combine results.



How to partition?

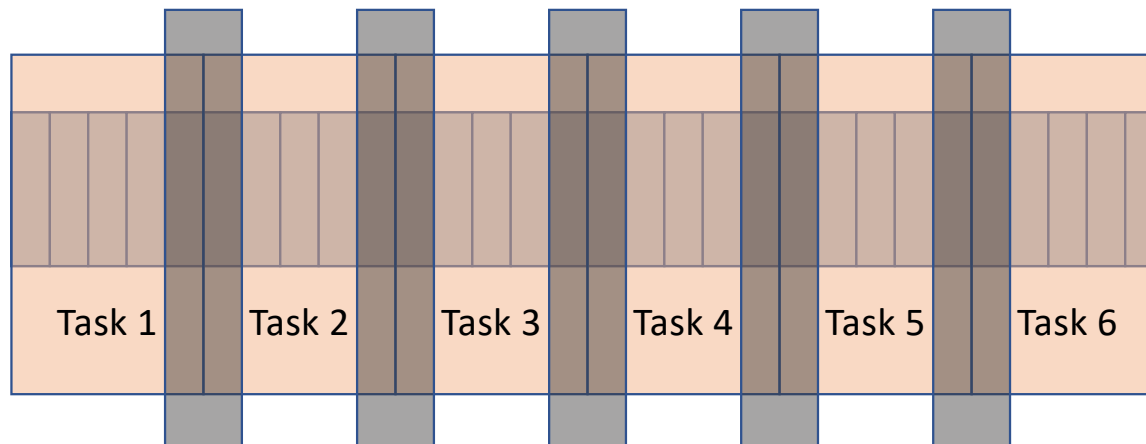
- How many parts
- How to divide
 - Partition 0, 0,4,8,12 ...
 - Partition 1,1,5,9,13 ...
 - Partition 2,2,6,10,14 ...
 - Partition 3,3,7,11,15 ...
- Or
 - Partition 0, 0 – 25 %
 - Partition 1, 25 – 50%
 - Partition 2, 50 – 75%
 - Partition 3, 75 – 100%

Array of data



Edge complications

- Each task runs on its own sub region, often they need to share edge data
- Update/Reading needs to be synchronised with neighbouring tasks



Parallel Loops

- Loops can be successfully parallelised IF
 - The order of execution is not important.
 - Each iteration is independent of each other.
 - Each Iteration has sufficient work to compensate for scheduling overhead.
- If insufficient work per loop iteration consider creating an inner loop.

Cost of loop parallelism

Time = $\frac{\text{Number of Tasks} * (\text{Task Scheduling Cost} + (\text{IterationCost} * \text{IterationsPerTask}))}{\text{Number of Cores}}$

Number of Cores

- Task Scheduling Cost = 50
- Iteration Cost = 1
- Total Iterations = 100,000
- Sequential Time would be 100,000

Number of Cores	Iterations Per Task	Number of Tasks	Time
2	100	1000	75000
2	50000	2	50050
4	100	1000	37500
4	25000	4	25050
256	100	1000	586
256	391	256	440

Functional programming ideal for parallelism

- Input => f() => f2() => f3() => output
- Immutable data

LINQ is a form of functional programming


- LINQ is readable
- PLINQ as readable, but in parallel
 - Free lunch
- Works with `IEnumerable<T>`
 - Better with `Array` and `List<T>`
 - Best when not order sensitive

PLINQ

- Embarrassingly parallel
- PLINQ not always faster
- LINQ is NOT a performance technology
 - Boosting it doesn't mean best
- Brings parallelism to the masses

Travelling salesman problem

- Calculate the shortest path, between a series of towns, starting and finishing at the same location
- Possible combinations n Factorial
- 15 City produces 1307674368000 combinations
- Took 8 Cores appx. 18 days to calculate exact shortest route.
- Estimate 256 Cores would take appx. 0.5 days




Monte Carlo

- Some algorithms even after heavy parallelisation are too slow.
- Often some answer in time T is often better than a perfect answer in $T+$
 - Real time speech processing
 - Route planning
- Machines can make educated guesses too..


Monte Carlo - Single Core

% Error for a 15 city route allowing 2 seconds for calculation using single core




Monte Carlo – 8 Cores

% Error for a 15 city route allowing 2 seconds for calculation using eight cores



Average % Error

Average % Error for a 15 city route allowing 2 seconds for calculation



Monte Carlo

- Each core runs independently making guesses.
 - Scales, more cores greater the chance of a perfect guess.
 - As we approach 256 cores...looks very attractive.
- Strategies
 - Pure Random
 - Hybrid, Random + Logic

Summary

- Lots of ways to parallelise
 - Frameworks makes things easier, but not trivial
 - PLINQ enables safe parallel programming to the masses
- Always benchmark
- For best performance tune to hardware